



DEEP LEARNING

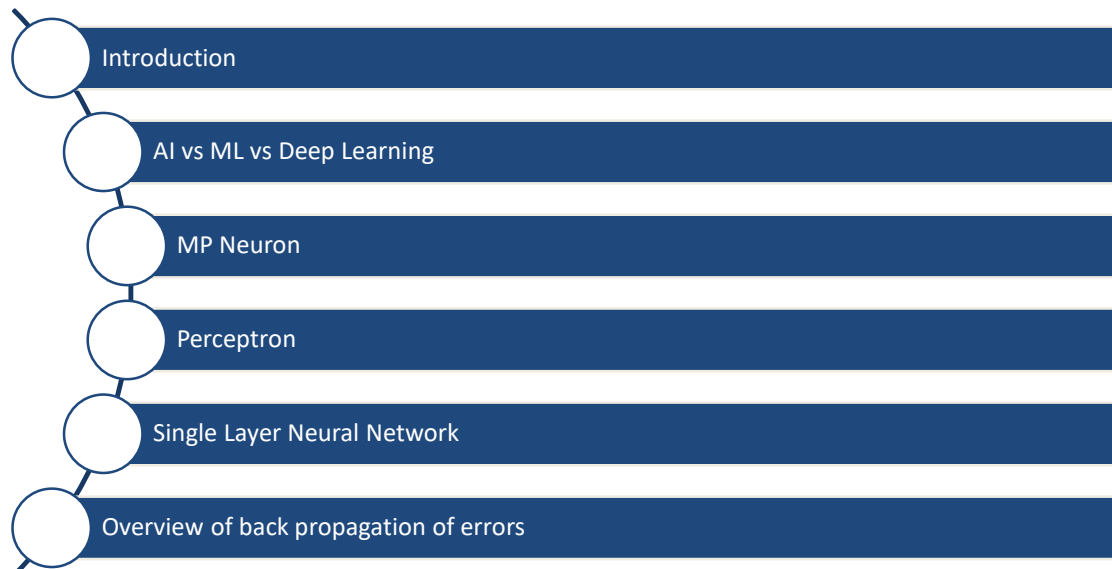
Introduction to Machine Learning, Artificial Intelligence and Deep Learning

Session 06

Pramod Sharma
pramod.sharma@prasami.com

2

Agenda



7/4/2023

pra-sâmi

3

References

- ❑ Deep Learning, Ian Goodfellow, Yoshua Bengio, Aaron Courville
- ❑ Neural Networks and Learning Machines, Simon Haykin
- ❑ Pattern Recognition and Machine Learning, Christopher M. Bishop
- ❑ Deep Learning with Python - François Chollet
- ❑ Hands-On Machine Learning with Scikit-Learn and TensorFlow
- ❑ TensorFlow Deep Learning Cookbook
- ❑ Reinforcement Learning with TensorFlow: A Beginner's Guide to Designing Self-learning Systems with TensorFlow and OpenAI Gym Sayon Dutta
- ❑ Hands-On Reinforcement Learning with Python: Master Reinforcement and Deep Reinforcement Learning Using OpenAI Gym and TensorFlow Sudharsan Ravichandiran
- ❑ Deep Reinforcement Learning Hands-On: Apply Modern RL Methods, with Deep Q-networks, Value Iteration, Policy Gradients, TRPO, AlphaGo Zero and More Maxim Lapan

7/4/2023

pra-sâmi

4

Agent In Uncertain Environment

- ❑ Agents don't have complete knowledge about the world.
- ❑ Agents need to make (informed) decisions given their uncertainty.
- ❑ It isn't enough to assume what the world is like.
 - ❖ Example: wearing a seat belt.
- ❑ An agent needs to reason about its uncertainty.
- ❑ When an agent takes an action under uncertainty, it is gambling \Rightarrow probability



7/4/2023

pra-sâmi

5

Overview

- ❑ Nature is a continuum where as math is discrete values
 - ❖ Old film based images were continuous painting of colors where as digital images are pixels
- ❑ Brain works differently than our mathematical computations
- ❑ Brain is highly complex, nonlinear and parallel computer
- ❑ Neural networks are supposed to be inspired from
- ❑ Highly generalized form, a Neural Network is a mathematical model that simulate manner in which brain performs a task



All models are wrong... some models are useful!

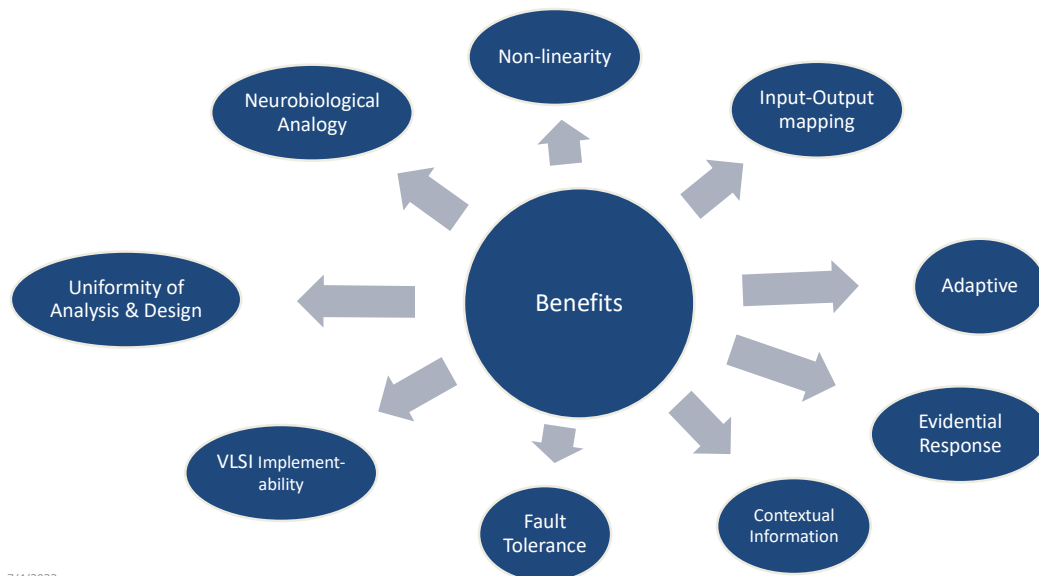
Is this how our brain works? Really!!

7/4/2023

pra-sâmi

6

Benefits of Neural Networks



7/4/2023

pra-sâmi

7

What has been achieved so far

- ❑ Learn to see and hear... so natural to humans but elusive to machines earlier
- ❑ Image classification
- ❑ Speech recognition
- ❑ Handwriting recognition
- ❑ Writing style recognition (who was the author)
- ❑ Improved machine translation
- ❑ Text-to-speech conversion
- ❑ Digital assistants such as Google Now and Amazon Alexa
- ❑ Little autonomous driving
- ❑ Improved ad targeting, as used by Google, Baidu, and Bing
- ❑ Ability to answer natural-language questions
- ❑ Superhuman games playing: chess, go...

Still long way to go...
Human-level general intelligence too far away...

7/4/2023

pra-sâmi

8

Deep learning is difficult!

7/4/2023

pra-sâmi

9

Neurons

7/4/2023

pra-sâmi

10

Neurons

- ❑ Take an example whether to go and play Cricket or not.
- ❑ Features:
 - ❖ Is it raining?
 - ❖ Is it too hot?
 - ❖ Have I completed my homework?
 - ❖ Are sufficient players ready?
 - ❖ Is cricket equipment ready?
 - ❖ Is ground available?
- ❑ Depending on the feature values, you may get to play or not
- ❑ Features like homework and availability of ground can be considered as 'inhibitory'.

id	Dry Weather	Low Temp.	Homework Done	Team Members	Equipment	Ground	Played
1	1	1	1	1	0	1	1
2	1	1	1	1	1	1	1
3	1	1	1	1	1	1	1
4	0	1	0	1	1	1	0
5	0	0	1	1	1	0	0
6	0	0	0	0	0	1	0

- ❑ Notes :
 - ❖ Aggregator function is sum and threshold can be 3.
 - ❖ Assign 0 or 1 if a parameter is in favor or not

Given sufficient data point, we can train an algorithm to make such simple decisions for us.

7/4/2023

pra-sâmi

11

MP Neuron

- ❑ In 1943 Warren S. McCulloch, a neuroscientist, and Walter Pitts, a logician, published "A logical calculus of the ideas immanent in nervous activity" in the Bulletin of Mathematical Biophysics
- ❑ In this paper McCulloch and Pitts tried to understand how the brain could produce highly complex patterns by using many basic cells that are connected together
- ❑ These basic brain cells are called neurons, and McCulloch and Pitts gave a highly simplified model of a neuron in their paper
- ❑ The McCulloch and Pitts model of a neuron, which we will call an MCP neuron for short, has made an important contribution to the development of artificial neural networks -- which model key features of biological neurons


7/4/2023

pra-sâmi

12

MP Neuron

- ❑ Neurons receive signals and produce a response
- ❑ In this model:
 - ❖ All inputs are binary i.e. [0,1]
 - ❖ Inputs are "inhibitory" or "excitatory".
 - ❖ Inhibitory have maximum influence on the model
 - ❖ It has an aggregator 'g' and a function 'f'
 - ❖ There is a threshold
 - ❖ If g is more than threshold, $\hat{y} = 1$ else 0



$$\hat{y} = 0 \text{ if any } x_i \text{ is inhibitory, else } g(x) = \sum x_i$$


$$\hat{y} = 1 \text{ if } g(x) \geq \text{threshold} \text{ else } \hat{y} = 0$$

7/4/2023

pra-sâmi

13

MP Neuron




id	Dry Weather	Low Temp	Homework Done	Team Members	Equipment	Ground	Sum	Played
1	1	1	1	1	0	1	5	1
2	1	1	1	1	1	1	6	1
3	1	1	1	1	1	1	6	1
4	0	1	0	1	1	1	4	0
5	0	0	1	1	1	0	3	0
6	0	0	0	0	0	1	1	0

7/4/2023

pra-sâmi

14

MP Neuron



id	Dry Weather	Low Temp	Homework Done	Team Members	Equipment	Ground	Sum	Played
1	1	1	1	1	0	1	5	1
2	1	1	1	1	1	1	6	1
3	1	1	1	1	1	1	6	1
4	0	1	0	1	1	1	4	0
5	0	0	1	1	1	0	3	0
6	0	0	0	0	0	1	1	0

The logic is straight forward. Let's implement this model on a dataset.


7/4/2023

pra-sâmi

15

Code Example1 – MP Neurons

- ❑ Need a dataset with plenty of features and binary output
- ❑ Use Breast Cancer dataset from scikit-learn
 - ❖ `data = sklearn.datasets.load_breast_cancer()`
- ❑ Its features are a continuous and we need binary
 - ❖ Use pandas `pd.cut` to bin the columns
 - ❖ `X_bin = X . apply (pd.cut, bins=2, labels=[1,0])`
 - ❖ For `b` in range `[0, num_features+1]`
 - Sum it by row and compare with `b`



7/4/2023

pra-sâmi

16


Perceptrons

7/4/2023

pra-sâmi

17

Solution to Equation of Perceptron



Ian Goodfellow
Yoshua Bengio
Aaron Courville

7/4/2023

pra-sâmi

18

To play or not to play...

id	Rains	Temp	Homework	Team Members	Equipment	Ground	Played
1	0	38	1	15	0	600	1
2	0	25	1	15	1	800	1
3	0	26	1	15	1	1000	1
4	5	27	1	10	1	600	0
5	20	23	0	8	1	1800	0
6	30	22	0	6	0	600	0

□ Features:

- ❖ Rains in millimeter
- ❖ Temperature in ° C
- ❖ Homework completed? – 0 : No; 1: Yes
- ❖ Team members : How many team members are ready to play?
- ❖ Is cricket equipment available?
- ❖ Ground: per hour rent in Rupees/hour

7/4/2023

pra-sâmi

19

Weights

- ❑ Each of the feature has different importance
- ❑ To assign importance to each of the feature, we use weights!
- ❑ Values of each features are in different order of magnitude
 - ❖ Summation is not going to work
 - ❖ Scale the features between 0 and 1

id	Rains	Temp	Homework	Team Members	Equipment	Ground	Played
1	0	38	1	15	0	600	1
2	0	25	1	15	1	800	1
3	0	26	1	15	1	1000	1
4	5	27	1	10	1	600	0
5	20	23	0	8	1	1800	0
6	30	22	0	6	0	600	0

- ❑ Note:
 - ❖ Variation in features have different bearing on the results
 - ❖ Team members → higher the better
 - ❖ Ground cost → lower the better

7/4/2023

pra-sâmi

20

Perceptron

- ❑ In MP Neuron Model,
 - ❖ All inputs had same weights
 - ❖ Threshold ' w_0 ' could take limited values
 - ❖ Every feature needed to be [0,1]
- ❑ Perceptron model introduced different weights to different inputs features
- ❑ Real values are also accepted
 - ❖ Temperatures are in tens and ground rent is in hundreds.
 - ❖ Min – Max – Scaler to compensate for huge difference in values
- ❑ Threshold ' w_0 ' can take any value
- ❑ Outputs are still [0, 1]

7/4/2023

pra-sâmi

21

Perceptron

□ Loss Function:

- ❖ A correction is applied on the outputs
- ❖ To adjust values of ' w_i ' to reach right results
- ❖ It would also give us indications of what weights to be fixed to arrive at the solution

□ Activation function $g(x)$ is applied as follows:

- ❖ If $\sum x_i \cdot w_i \geq w_0 \Rightarrow \hat{y} = 1$
- ❖ If $\sum x_i \cdot w_i < w_0 \Rightarrow \hat{y} = 0$

7/4/2023

pra-sâmi

22

Perceptron – Data Preprocessing

- Lets consider “Ground” and “Team Members” as features and its associated weights to arrive at the solution.

id	Rains	Temp	Homework	Team Members	Equipment	Ground	Played
1	0	38	1	15	0	600	1
2	0	25	1	15	1	800	1
3	0	26	1	15	1	1000	1
4	5	27	1	10	1	600	0
5	20	23	0	8	1	1800	0
6	30	22	0	6	0	600	0

7/4/2023

pra-sâmi

23

Perceptron – Data Preprocessing

- Scaled Data (all columns to be between 0 and 1)

id	Rains	Temp	Homework	Team Members	Equipment	Ground	Played
1	0.00	0.00	1.00	1.00	0.00	1.00	1
2	0.00	0.81	1.00	1.00	1.00	0.83	1
3	0.00	0.75	1.00	1.00	1.00	0.67	1
4	-0.17	0.69	1.00	0.44	1.00	1.00	0
5	-0.67	0.94	0.00	0.22	1.00	0.00	0
6	-1.00	1.00	0.00	0.00	0.00	1.00	0

- What about reverse correlation
- Two option to address reverse correlation
 - ❖ Take negative of values
 - ❖ Use negative weight

7/4/2023

pra-sâmi

24

Perceptron – Weights

- Weights – consider importance of each of the feature

id	Threshold	Team Members		Ground		Calculations	Likely	Played	Loss
	w0	x1	w1	x2	w2	$w0+x1*w1+x2*w2$	(y_hat)	(y)	(y-y_hat)^2
1	-1.00	1.00	1.10	1.00	1.00	1.10	1	1	0
2	-1.00	1.00	1.10	0.83	1.00	0.93	1	1	0
3	-1.00	1.00	1.10	0.67	1.00	0.77	1	1	0
4	-1.00	0.44	1.10	1.00	1.00	0.49	1	0	1
5	-1.00	0.22	1.10	0.00	1.00	-0.76	0	0	0
6	-1.00	0.00	1.10	1.00	1.00	0.00	1	0	1

7/4/2023

pra-sâmi

25

Perceptron – Weights and Loss

- ❑ Our best solution would be where ground truth and predicted values are same
- ❑ Loss is some function of ground truth and predicted values
- ❑ And we want it to be cumulative, Square of difference looks promising
 - ❖ $\ell(\hat{y}, y) = (y - \hat{y})^2$
 - ❖ Our overall loss was 2.
- ❑ By adjusting weights (w_1, w_2) and threshold (w_0) we can bring the loss to minimum (zero in this case)

id	Threshold	Team Members		Ground		Calculations	Likely	Played	Loss
	w0	x1	w1	x2	w2	$w_0 + x_1 * w_1 + x_2 * w_2$	(y_hat)	(y)	$(y - y_hat)^2$
1	-2.00	1.00	1.10	1.00	1.40	0.50	1	1	0
2	-2.00	1.00	1.10	0.83	1.40	0.27	1	1	0
3	-2.00	1.00	1.10	0.67	1.40	0.03	1	1	0
4	-2.00	0.44	1.10	1.00	1.40	-0.11	0	0	0
5	-2.00	0.22	1.10	0.00	1.40	-1.76	0	0	0
6	-2.00	0.00	1.10	1.00	1.40	-0.60	0	0	0

7/4/2023

pra-sâmi

26

Perceptron – Weights and Loss

- ❑ Our best solution would be where ground truth and predicted values are same.
- ❑ Loss is some function of ground truth and predicted values
- ❑ And we want it to be cumulative, Square of difference looks promising
 - ❖ Square of difference looks promising
 - ❖ $\ell(\hat{y}, y) = (y - \hat{y})^2$
- ❑ Our overall loss was 2.
- ❑ By adjusting weights (w_1, w_2) and threshold (w_0) we can bring the loss to minimum (zero in this case)

id	Threshold	Team Members		Ground		Calculations	Likely	Played	Loss
	w0	x1	w1	x2	w2	$w_0 + x_1 * w_1 + x_2 * w_2$	(y_hat)	(y)	$(y - y_hat)^2$
1	-2.00	1.00	1.10	1.00	1.40	0.50	1	1	0
2	-2.00	1.00	1.10	0.83	1.40	0.27	1	1	0
3	-2.00	1.00	1.10	0.67	1.40	0.03	1	1	0
4	-2.00	0.44	1.10	1.00	1.40	-0.11	0	0	0
5	-2.00	0.22	1.10	0.00	1.40	-1.76	0	0	0
6	-2.00	0.00	1.10	1.00	1.40	-0.60	0	0	0

Voilà!


We made it!!

7/4/2023


pra-sâmi

27

Perceptron



- We can represent : $g = w_0 + x_1 * w_1 + x_2 * w_2$
 - ✦ As $g = [x_1, x_2] \cdot \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} + w_0$
- Given: $W = \begin{bmatrix} 1.1 \\ 1.4 \end{bmatrix}$ and $w_0 = -2$
 - ✦ $g = [x_1, x_2] \cdot \begin{bmatrix} 1.1 \\ 1.4 \end{bmatrix} - 2$
 - ✦ $g = 1.1 * x_1 + 1.4 * x_2 - 2$



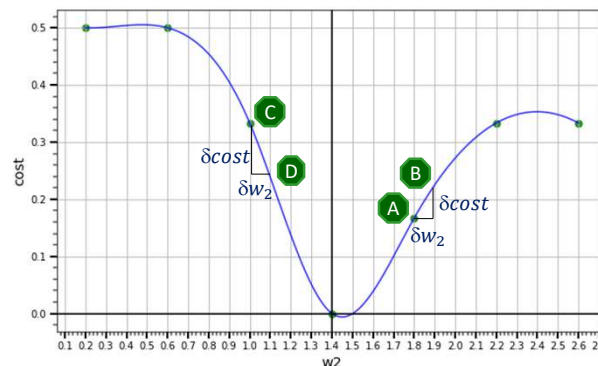
7/4/2023

pra-sâmi

28

Perceptron – Gradient Descent

- w_0, w_1, w_2 need to be adjusted to arrive at most optimal solution i.e. lowest point on the graph.
- Assume that w_0 is fixed at -2, and w_1 at 1.1 and w_2 varies from 0 to 3 (only one variable considered to make plotting simple)
- From point A to B, slope is positive hence w_2 value needs to be decreased
- From point C to D slope is negative hence w_2 needs to be increased.



7/4/2023

pra-sâmi

29


Perceptron – Activation Function

- So we based our entire calculations on:

$$z = w_0 + x_1 * w_1 + x_2 * w_2$$



But that's an equation of straight line! 😊
What happened to all those 'inhibitory' features?





7/4/2023

pra-sâmi

30

Non Linear Activation function




Non-linear Activation function

7/4/2023

pra-sâmi

31

Perceptron with non-linear activation function



□ Given:

$$\diamond W = \begin{bmatrix} 2 \\ -3 \end{bmatrix} \text{ and } b = 1$$

$$\diamond \hat{y} = \sigma \left([x_1, x_2] \cdot \begin{bmatrix} 2 \\ -3 \end{bmatrix} + 1 \right)$$

$$\diamond \hat{y} = \sigma \left(\underbrace{1 + 2 * x_1 - 3 * x_2}_z \right)$$

$$\square \hat{y} = \sigma(z);$$

□ Lets use sigmoid function for σ .

$$\diamond \hat{y} = \frac{1}{(1+e^{-z})}$$

7/4/2023

pra-sâmi

32

Perceptron with non-linear activation function

$$\square \hat{y} = \sigma(1 + 2 * x_1 - 3 * x_2)$$

□ For $X = [-3, 4]$

$$\diamond \hat{y} = \sigma(1 + 2 * (-3) - 3 * 4)$$

$$\diamond \hat{y} = \sigma(1 - 6 - 12)$$

$$\diamond \hat{y} = \sigma(-17)$$

$$\diamond \hat{y} = 0.0$$


□ Similarly, for $X = [1, -2]$

$$\diamond \hat{y} = \sigma(1 + 2 * 1 - 3 * (-2))$$

$$\diamond \hat{y} = \sigma(1 + 2 - 6)$$


$$\diamond \hat{y} = \sigma(9)$$

$$\diamond \hat{y} = 1.0$$



7/4/2023

pra-sâmi



35

Exercise 2 : Computational Graph

- Given a Cost Function J
 - ❖ $J(w, x, b) = 3 * (b + x * w)$
- Calculate $\frac{\partial J}{\partial w}$, $\frac{\partial J}{\partial x}$ and $\frac{\partial J}{\partial b}$
- Calculate slope at point :
 - ❖ $b = 6$
 - ❖ $w = 3$
 - ❖ $x = 2$



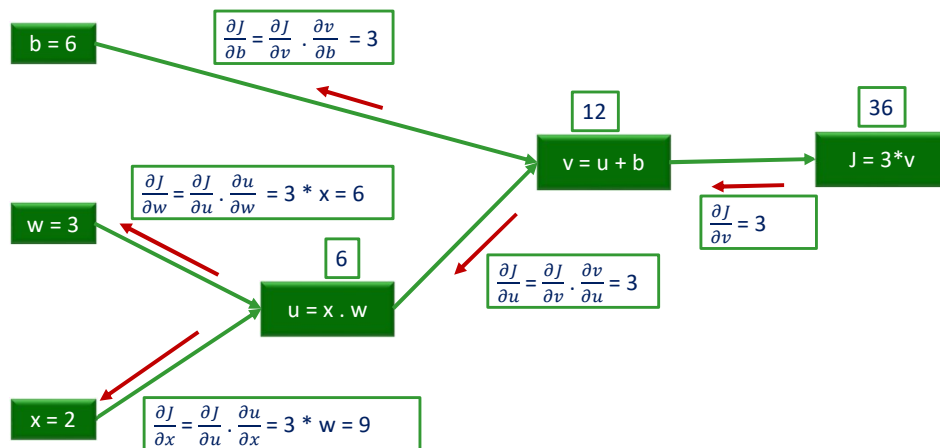
7/4/2023

pra-sâmi

36

Exercise - Solution

- Given a Cost Function $J(w, x, b) = 3 * (b + w * x)$




7/4/2023

pra-sâmi

37

Consider Single Path... MLE

- ❑ Maximum likelihood estimation, or MLE, is a framework for inference for finding the best statistical estimates of parameters from historical training data
 - ❖ Exactly what we are trying to do with the neural network
- ❑ In Classification, output is probability of it belonging to a class
 - ❖ Maximum likelihood estimation, seeks a set of model weights that minimize the difference between the predicted probability distribution and the Ground Truth [cross-entropy]
- ❑ In Regression problems:
 - ❖ Use the mean squared error (MSE) loss function or equivalent.




7/4/2023

pra-sâmi

38

Consider Single Path... Loss Function

- ❑ A function used to evaluate a candidate solution
- ❑ Helps to maximize or minimize the objective function
- ❑ Estimates how closely the distribution of predictions made by a model matches the ground truth (maximum likelihood)
- ❑ Under maximum likelihood framework, the error between two probability distributions is measured using cross-entropy
 - ❖ Hence $\ell(\hat{y}, y) = -[y * \log(\hat{y}) + (1 - y) * \log(1 - \hat{y})]$



7/4/2023

pra-sâmi

39

Cost Function

$$\hat{y} = \sigma(\sum W * X + b)$$

$$\text{Where } \sigma(z) = \frac{1}{1+e^{-z}}$$

Loss function:

- ❖ A parameter which defines how good our outputs are i.e.
- ❖ How far our predicted values ' \hat{y} ' (y hat) were from ground truth ' y '

For logistic regression

- ❖ $\text{Loss}(\hat{y}, y) = -(y \cdot \log \hat{y} + (1-y) \cdot \log(1-\hat{y}))$
- ❖ Loss function is for an instance
- ❖ In case of binary classification, $\text{Loss}(\hat{y}, y) = -y \cdot \log \hat{y}$

Cost Function: Its a sum of losses for all instances

$$J(W, b) = \frac{1}{m} (\sum \text{Loss}(\hat{y}, y))$$

$$= -\frac{1}{m} (\sum (y \cdot \log \hat{y} + (1-y) \cdot \log(1-\hat{y})))$$

For binary classification:

$$J(W, b) = \frac{1}{m} (\sum \text{Loss}(\hat{y}, y))$$

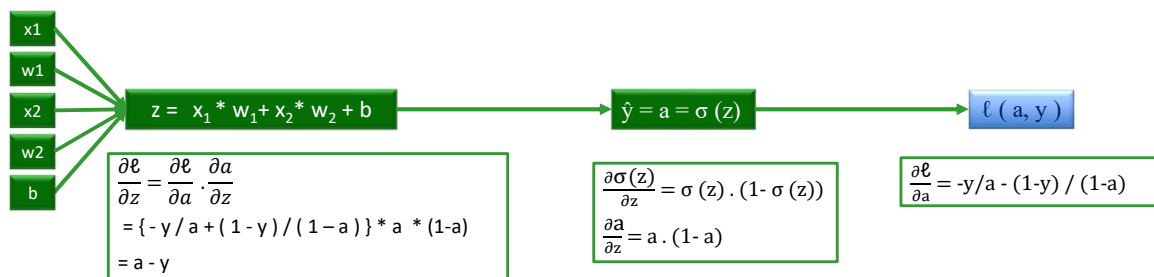
$$= -\frac{1}{m} (\sum (y \cdot \log \hat{y}))$$

7/4/2023

pra-sâmi

40

Forward and Back Propagation



$$z = X * W + b$$

$$\hat{y} = a = \sigma(z)$$

$$\sigma(z) = \frac{1}{1+e^{-z}}$$

$$\ell(a, y) = -[y * \log(a) + (1-y) * \log(1-a)]$$

For binary classification:

$$\ell(a, y) = -y * \log(a)$$



$$\frac{\partial \ell}{\partial w_1} = x_1 \cdot \frac{\partial \ell}{\partial z} = x_1 \cdot (a-y)$$

$$\frac{\partial \ell}{\partial w_2} = x_2 \cdot \frac{\partial \ell}{\partial z} = x_2 \cdot (a-y)$$

$$\frac{\partial \ell}{\partial b} = \frac{\partial \ell}{\partial z} = (a-y)$$



$$w_1 = w_1 - \alpha * \frac{\partial \ell}{\partial w_1} = w_1 - \alpha * x_1 * (a-y)$$

$$w_2 = w_2 - \alpha * \frac{\partial \ell}{\partial w_2} = w_2 - \alpha * x_2 * (a-y)$$

$$b = b - \alpha * \frac{\partial \ell}{\partial b} = b - \alpha * (a-y)$$

Where α is learning rate. The cost function is

$$J(W, b) = \frac{1}{m} (\sum \ell(a, y))$$

$$\text{Hence } \frac{\partial J}{\partial w_1} = \frac{1}{m} * (\sum \frac{\partial \ell(a, y)}{\partial w_1})$$

7/4/2023

pra-sâmi

41

So where are the hidden layers!!!


7/4/2023

pra-sâmi

42

Hidden Layers

id	Threshold	Team Members		Ground	
	x0	x1	w1	x2	w2
1	-2.00	1.00	1.10	1.00	1.40
2	-2.00	1.00	1.10	0.83	1.40
3	-2.00	1.00	1.10	0.67	1.40
4	-2.00	0.44	1.10	1.00	1.40
5	-2.00	0.22	1.10	0.00	1.40
6	-2.00	0.00	1.10	1.00	1.40




7/4/2023

pra-sâmi

43

Hidden Layers

id	Rains	Temp	Homework	Team Members	Equipment	Ground	Played
1	0.00	0.00	1.00	1.00	0.00	1.00	1
2	0.00	0.81	1.00	1.00	1.00	0.83	1
3	0.00	0.75	1.00	1.00	1.00	0.67	1
4	-0.17	0.69	1.00	0.44	1.00	1.00	0
5	-0.67	0.94	0.00	0.22	1.00	0.00	0
6	-1.00	1.00	0.00	0.00	0.00	1.00	0




7/4/2023

pra-sâmi

44

Hidden Layers

id	Rains	Temp	Homework	Team Members	Equipment	Ground	Played
1	0.00	0.00	1.00	1.00	0.00	1.00	1
2	0.00	0.81	1.00	1.00	1.00	0.83	1
3	0.00	0.75	1.00	1.00	1.00	0.67	1
4	-0.17	0.69	1.00	0.44	1.00	1.00	0
5	-0.67	0.94	0.00	0.22	1.00	0.00	0
6	-1.00	1.00	0.00	0.00	0.00	1.00	0




7/4/2023

pra-sâmi

45

Hidden Layers

id	Rains	Temp	Homework	Team Members	Equipment	Ground	Played
1	0.00	0.00	1.00	1.00	0.00	1.00	1
2	0.00	0.81	1.00	1.00	1.00	0.83	1
3	0.00	0.75	1.00	1.00	1.00	0.67	1
4	-0.17	0.69	1.00	0.44	1.00	1.00	0
5	-0.67	0.94	0.00	0.22	1.00	0.00	0
6	-1.00	1.00	0.00	0.00	0.00	1.00	0

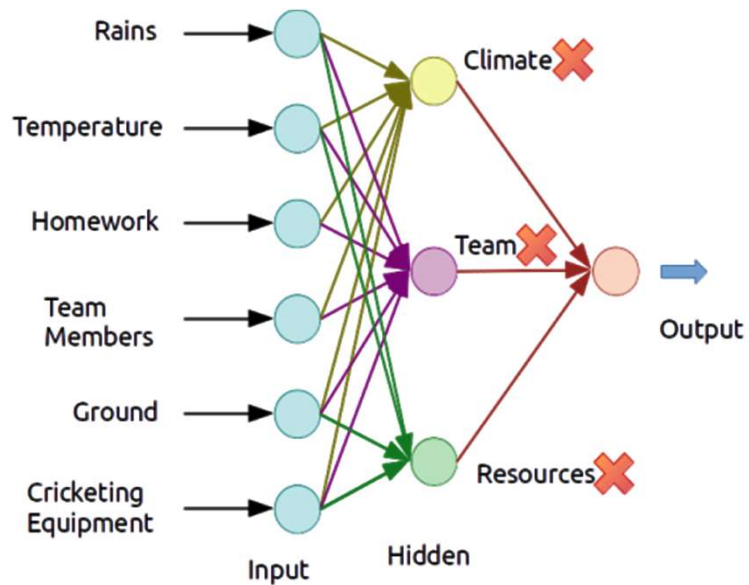


7/4/2023

pra-sâmi

46

Hidden Layers



7/4/2023

pra-sâmi


47

Hidden Layers



48

Two Major Conventions

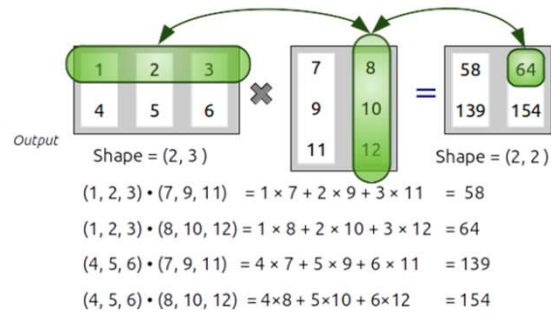



7/4/2023

pra-sâmi

49

Two Major Conventions






7/4/2023


pra-sâmi

50

Two Major Conventions



7/4/2023



pra-sâmi

51 Two M

Voilà !! Its time to update the resume!

(m, 1) pra-sâmi

7/4/2023

52 Reflect...

- ❑ How many type of layers Deep Learning Algorithms have?
 - a. 2
 - b. 3
 - c. 4
 - d. 5
- ❑ Answer : b
- ❑ The first layer is called the?
 - a. Input Layer
 - b. Output Layer
 - c. Hidden Layer
 - d. None of The Above
- ❑ Answer : a
- ❑ Which of the following is/are Limitations of deep learning?
 - a. Data labeling
 - b. Obtain huge training datasets
 - c. Both A and B
 - d. None of the above
- ❑ Answer : c
- ❑ Deep learning algorithms are _____ more accurate than machine learning algorithm in image classification.
 - a. 33%
 - b. 37%
 - c. 40%
 - d. 41%
- ❑ Answer : d

7/4/2023

pra-sâmi

53

Reflect...

- ☐ In which of the following applications can we use deep learning to solve the problem
 - a. Protein structure prediction
 - b. Prediction of chemical reactions
 - c. Detection of exotic particles
 - d. All of the above
- ☐ Answer : d

- ☐ The number of nodes in the input layer is 10 and the hidden layer is 5. The maximum number of connections from the input layer to the hidden layer are:
 - a. 50
 - b. less than 50
 - c. more than 50
 - d. It is an arbitrary value
- ☐ Answer : a

7/4/2023

pra-sâmi

54

Next Session - Coding Perceptron Model in Python

7/4/2023

pra-sâmi

55



7/4/2023

pra-sâmi


EXTRA MATERIAL

pra-sâmi

57

Applications

- The properties of neural networks define where they are useful
- Typical Network
 - ❖ Can learn complex mappings from inputs to outputs, based solely on samples
 - ❖ Difficult to analyse
 - ❖ Firm predictions about neural network behaviour difficult;
 - Unsuitable for safety-critical applications.
 - ❖ Require limited understanding from trainer, who can be guided by heuristics




7/4/2023

pra-sâmi

58

Applications

- The properties of neural networks define where they are useful
- Typical Network
 - ❖ Can learn complex mappings from inputs to outputs, based solely on samples
 - ❖ Difficult to analyse
 - ❖ Firm predictions about neural network behaviour difficult;
 - Unsuitable for safety-critical applications.
 - ❖ Require limited understanding from trainer, who can be guided by heuristics



7/4/2023


pra-sâmi

59

Applications

□ Stock market prediction

- ❖ “Technical trading” refers to trading based solely on known statistical parameters; e.g. previous price
- ❖ Neural networks have been used to attempt to predict changes in prices.
- ❖ Difficult to assess success or otherwise
 - Since companies using these techniques are reluctant to disclose information.



□ Mortgage assessment

- ❖ Assess risk of lending to an individual
- ❖ Difficult to decide on marginal cases
- ❖ Neural networks have been trained to make decisions, based upon the opinions of expert underwriters
- ❖ Neural network produced a 12% reduction in delinquencies compared with human experts



7/4/2023

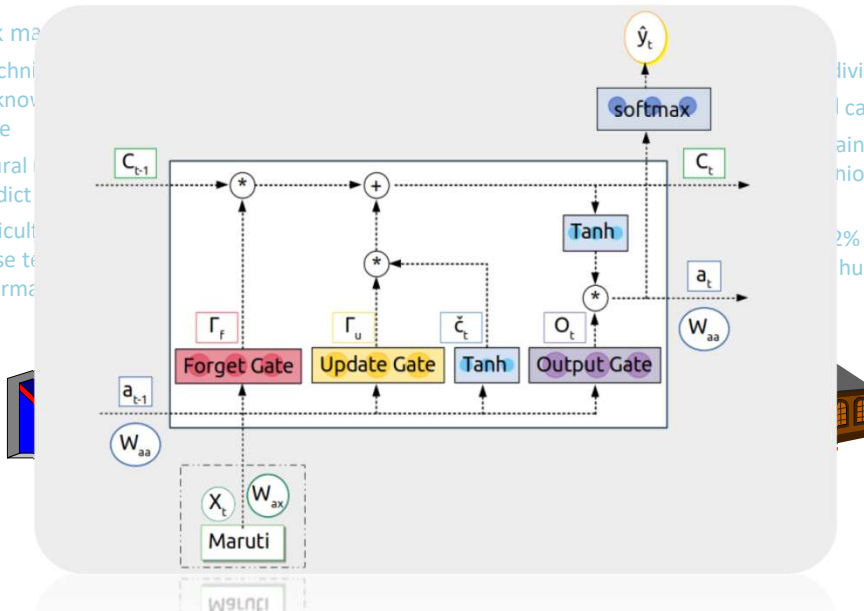
pra-sâmi

60

Applications

□ Stock market prediction

- ❖ “Technical trading” refers to trading based solely on known statistical parameters; e.g. previous price
- ❖ Neural networks have been used to attempt to predict changes in prices.
- ❖ Difficult to assess success or otherwise
 - Since companies using these techniques are reluctant to disclose information.



7/4/2023

pra-sâmi

61

Applications

- ALVINN: Autonomous Land Vehicle In a Neural Network





Figure 1: ALVINN Architecture

7/4/2023

- Sonar target recognition**

- ❖ Distinguish mines from rocks on sea-bed
- ❖ The neural network is provided with a large number of parameters which are extracted from the sonar signal.
- ❖ The training set consists of sets of signals from rocks and mines.

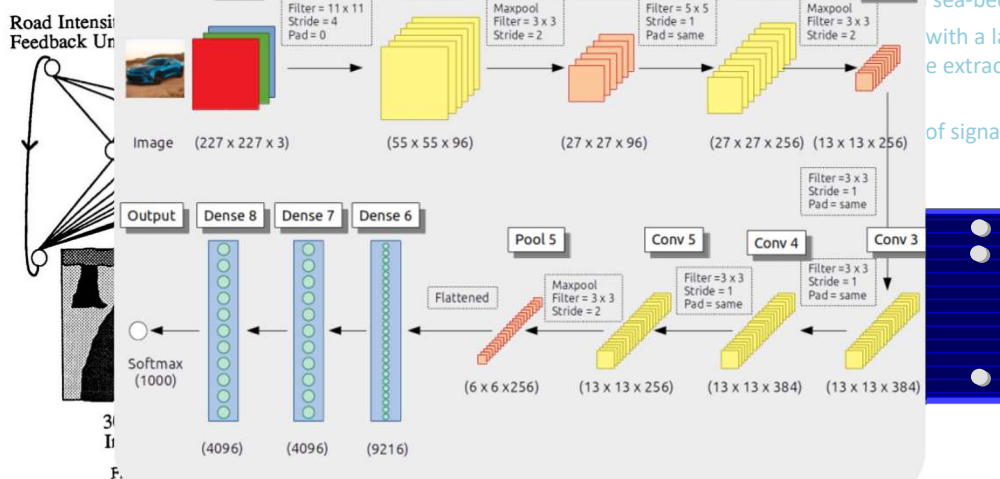


pra-sâmi

62

Applications

- ALVINN: Autonomous Land Vehicle In a Neural Network



7/4/2023

pra-sâmi

63

Applications

□ Engine management

- ❖ The behavior of a car engine is influenced by a large number of parameters
 - temperature at various points
 - fuel/air mixture
 - lubricant viscosity.
- ❖ Major companies have used neural networks to dynamically tune an engine depending on current settings



7/4/2023

□ Signature recognition

- ❖ Each person's signature is different.
- ❖ There are structural similarities which are difficult to quantify.
- ❖ Recognizes signatures to a high level of accuracy.
- ❖ Considers speed in addition to gross shape
- ❖ Makes forgery even more difficult.


pra-sâmi

64

Applications

□ Engine management


- ❖ The behavior of a car engine is influenced by a



7/4/2023

□ Signature recognition

- ❖ Each person's signature is different.



pra-sâmi

65

Derivation of Sigmoid

$$\begin{aligned}
 \partial a &= \partial \sigma(z) \\
 &= \frac{\partial}{\partial z} \left[\frac{1}{1 + e^{-z}} \right] \\
 &= \frac{\partial}{\partial z} (1 + e^{-z})^{-1} \\
 &= -(1 + e^{-z})^{-2} (-e^{-z}) \\
 &= \frac{e^{-z}}{(1 + e^{-z})^2} \\
 &= \frac{1}{1 + e^{-z}} \circ \frac{e^{-z}}{1 + e^{-z}} \\
 &= \frac{1}{1 + e^{-z}} \circ \frac{(1 + e^{-z}) - 1}{1 + e^{-z}} \\
 &= \frac{1}{1 + e^{-z}} \circ \left[\frac{1 + e^{-z}}{1 + e^{-z}} - \frac{1}{1 + e^{-z}} \right] \\
 &= \frac{1}{1 + e^{-z}} \circ \left[1 - \frac{1}{1 + e^{-z}} \right] \\
 &= \sigma(z) \circ (1 - \sigma(z)) \\
 &= a \circ (1 - a)
 \end{aligned}$$


7/4/2023

pra-sâmi

66


Biological Neural Nets

- ❑ Pigeons as art experts (Watanabe et al. 1995)
- ❑ Experiment:
 - ❖ Pigeon in Skinner box
 - ❖ Present paintings of two different artists (e.g. Chagall / Van Gogh)
 - ❖ Reward for pecking when presented a particular artist (e.g. Van Gogh)



7/4/2023

pra-sâmi



69

Biological Neural Nets

- ❑ Pigeons were able to discriminate between Van Gogh and Chagall
 - ❖ With 95% accuracy on train set (when presented with pictures they had been trained on)
 - ❖ Discrimination, still 85% successful for previously unseen paintings of the artists
- ❑ Pigeons do not simply memorise the pictures
 - ❑ They can extract and recognise patterns (the 'style')
 - ❑ They generalise from the already seen to make predictions
- ❑ This is what neural networks (biological and artificial) are good at (unlike conventional computer)

7/4/2023

pra-sâmi

70

Brain and Machine

- ❑ The Brain
 - ❖ Pattern Recognition
 - ❖ Association
 - ❖ Complexity
 - ❖ Noise Tolerance




- ❑ The Machine
 - ❖ Calculation
 - ❖ Precision
 - ❖ Logic

7/4/2023

pra-sâmi


71

The contrast in architecture



- ❑ The Von Neumann architecture uses a single processing unit;
 - ❖ Tens of millions of operations per second
 - ❖ Absolute arithmetic precision

- ❑ The brain uses many slow unreliable processors acting in parallel



7/4/2023

pra-sâmi

72

The biological inspiration

- ❑ Features of the Brain
 - ❖ Ten billion (10^{10}) neurons
 - ❖ On average, several thousand connections
 - ❖ Hundreds of operations per second
 - ❖ Die off frequently (never replaced)
 - ❖ Compensates for problems by massive parallelism
- ❑ The brain has been extensively studied by scientists
- ❑ Vast complexity prevents all but rudimentary understanding
- ❑ Even the behavior of an individual neuron is extremely complex
- ❑ Single “percepts” distributed among many neurons
- ❑ Localized parts of the brain are responsible for certain well-defined functions (e.g. vision, motion).



7/4/2023

pra-sâmi