



REGULARIZATIONS

Introduction to Machine Learning, Artificial Intelligence and Deep Learning

Session 11

Pramod Sharma
pramod.sharma@prasami.com

2

Agenda

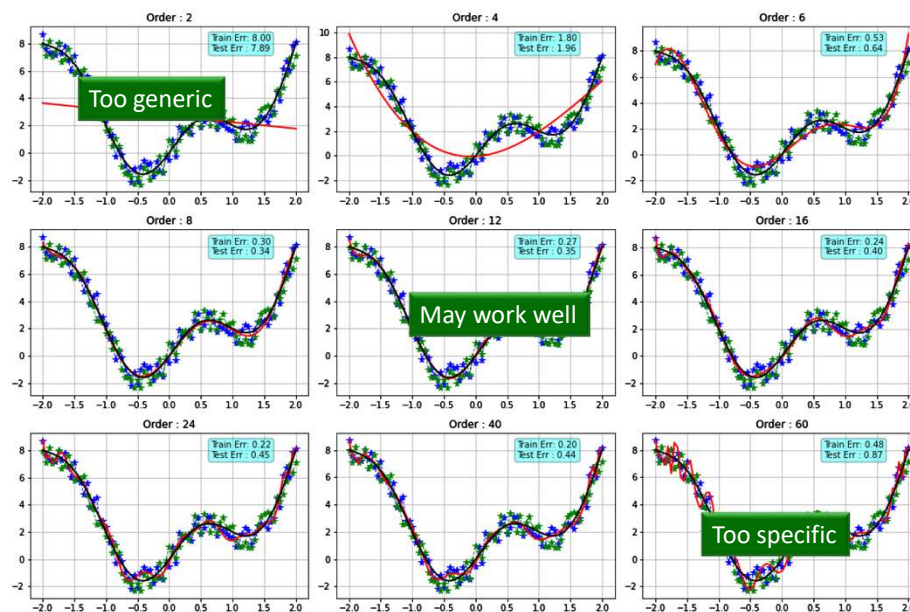


7/10/2023

pra-sâmi

3

Under-fitting vs. Over-fitting



7/10/2023


pra-sâmi

4

Regularization

- ❑ Regularization helps in avoiding over fitting by penalizing the coefficients
- ❑ In deep learning, it actually penalizes the weight matrices of the nodes
- ❑ Different Regularization Techniques in Deep Learning
 - ❖ L1 regularization
 - ❖ L2 regularizations
 - ❖ Dropouts
 - ❖ Early stopping
 - ❖ Data Augmentation

Most Libraries have tunable hyper-parameters!



7/10/2023

pra-sâmi

5

Weights vs. Bias


- ❑ For neural networks, we typically choose to use a parameter norm penalty Ω that penalizes only the weights at each layer and leaves the biases un-regularized.
- ❑ The biases typically require less data to fit accurately than the weights.
- ❑ Fitting the weight will requires observing both layer in a variety of conditions.
- ❑ Each bias controls only a single layer.
- ❑ This means that we do not induce too much variance by leaving the biases un-regularized.
- ❑ Also, regularizing the bias parameters can introduce a significant amount of under fitting.

7/10/2023

pra-sâmi

6

Effect of L2 Regularization



7/10/2023

pra-sâmi

7

Theory – Logistic Regression – L1 & L2

- Idea is to minimize Cost Function

$$\begin{aligned} \diamond J(W, b) &= \frac{1}{m} * \sum \ell(a, y) \\ \diamond &= -\frac{1}{m} \{y * \log(a) + (1-y) * \log(1-a)\} \end{aligned}$$

- A term is added to Cost function $\frac{\lambda}{2 * m} \cdot \|W\|_2^2$

$$\square J(W, b) = \frac{1}{m} * \sum \ell(a, y) + \frac{\lambda}{2 * m} \cdot \|W\|_2^2$$

7/10/2023

pra-sâmi

8

Theory – Logistic Regression – L1 & L2

$$\square J(W, b) = \frac{1}{m} * \sum \ell(a, y) + \frac{\lambda}{2 * m} \cdot \|W\|_2^2 + \frac{\lambda}{2 * m} \cdot b^2$$

- ✦ This is referred as L2 regularization
- ✦ **Regularization hyperparameter λ :** It is another parameter we tune...

$$\square \|W\|_2^2 = \sum_{j=1}^n w_j^2 = W^T \cdot W$$

- Here, we are using Euclidean Norm or L2 Norm
- Compared to W, bias b has fewer dimensions, hence, it is generally not considered
- If you add for b, $(\frac{\lambda}{2 * m} \cdot b^2)$... that's ok too
 - ✦ Although its effect will be minimal,
 - ✦ Better to leave it alone.

7/10/2023

pra-sâmi

9

Theory – Logistic Regression – L1 & L2

- Sometimes L1 too is used
- $J(W, b) = \frac{1}{m} * (\sum \ell(a, y)) + \frac{\lambda}{2 * m} \cdot \|W\|_1$
- Differentiation of $\frac{\lambda}{2 * m} \cdot \|W\|_1 = \frac{\lambda}{2 * m} \text{sign}(W)$
 - ❖ Will be infinitely small and will have insignificant impact on gradient descent

7/10/2023

pra-sâmi

10

Neural Network – Frobenius Norm

- In neural network, we have different layers with different weights
- So we look at its cumulative effect over all layers
- Hence the Cost function
 - ❖ $J(W, b) = J(W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}, W^{[3]}, b^{[3]} \dots)$
 - ❖ $J(W, b) = \frac{1}{m} * (\sum \ell(a, y)) + \frac{\lambda}{2 * m} \cdot \sum_{l=1}^L \sum (w_{i,j})^2$
 - ❖ $J(W, b) = \frac{1}{m} * \sum \{y * \log(a) + (1-y) * \log(1-a)\} + \frac{\lambda}{2 * m} \cdot \sum_{l=1}^L \|W^{[l]}\|^2$
 - ❖ Where $\|W^{[l]}\|^2_F = \sum_{i=1}^{n^{[l-1]}} \sum_{j=1}^{n^{[l]}} (w_{ij}^{[l]})^2$
 - W is $(n^{[l-1]}, n^{[l]})$ dimensional matrix
- It is called *Frobenius norm* of a matrix
- Also the Euclidean norm defined as the square root of the sum of the absolute squares of its elements

7/10/2023

pra-sâmi

11

Frobenius Norm of a Vector

$$\square \|A\|_F = \sqrt{\sum (a_{ij})^2}$$

i.e.

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix} = \sqrt{(1^2 + 2^2 + 3^2 + 4^2 + 5^2 + 6^2 + 7^2 + 8^2 + 9^2)}$$

7/10/2023

pra-sâmi

12

Updates to weights

□ Earlier

$$\diamond \partial W^{[l]} = X \cdot \partial z$$

$$\diamond \text{ And } W^{[l]} = W^{[l]} - \alpha \cdot \partial W^{[l]}$$

□ For Regularization we add an extra term at the end

$$\diamond \partial W^{[l]} = X \cdot \partial z + \frac{\lambda}{m} \cdot W^{[l]}$$

Mathematically, we can show that it is still a valid definition of $\partial W^{[l]}$

$$\diamond W^{[l]} = W^{[l]} - \alpha \cdot [X \cdot \partial z + \frac{\lambda}{m} \cdot W^{[l]}]$$

$$\diamond W^{[l]} = \left(1 - \frac{\alpha \cdot \lambda}{m}\right) W^{[l]} - \alpha \cdot X \cdot \partial z$$

Weight Decay


7/10/2023

pra-sâmi


13

Regularization : Dropout

Iteration 1



Iteration 2




7/10/2023

pra-sâmi


14

Regularization : Dropout

Original



With Dropout



7/10/2023

pra-sâmi

15

Regularization : Early Stopping

- ❑ How long to train the model?
- ❑ Duration of training → under – fit or over – fit
- ❑ Train the model to the point where its performance on test set is best!
- ❑ Very simple and very effective

How:

- ❑ Train the model and monitor performance
- ❑ Save weight every time the performance improves
- ❑ Stop training if performance has not improved for N epochs
- ❑ It's the last parameter to tune
 - ❖ Repeated early stopping may lead to over-fitting the validation set
 - ❖ Example : K-fold

7/10/2023

pra-sâmi

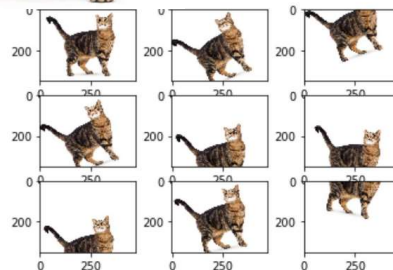
16

Regularization : Data Augmentation

- ❑ Where limited data is available for training the model (when is it not!)
- ❑ Very effective in image identification
- ❑ Most libraries have Image Generators (parameter driven)
 - ❖ Horizontal and Vertical Shift
 - ❖ Horizontal and Vertical Flip
 - ❖ Random Rotation
 - ❖ Random Brightness / Contrast
 - ❖ Random Zoom
 - ❖ Random Noise



<https://towardsdatascience.com/image-augmentation-for-deep-learning-histogram-equalization-a71387f609b2>



7/10/2023

pra-sâmi

17



7/10/2023

pra-sâmi